




# Square and Cube Root Optimization



Team 25:  
Keegan Kavanagh  
Jon Squire



# Tasks

---

1. Compute the calculation of square roots and cube roots
2. Optimize the resulting algorithm of for the ARM machine
3. Analyze how different optimizations affect the square root and cube root function differently

# Requirements

---

- Compiled using gcc version 4.3.2 for the 2.6.29.4-FriendlyARM platform
- Written in C
- Square root input range of  $[1,4)$
- Cube root input range of  $[1,8)$
- 32 bit output

# The CCM Algorithm for Computing Roots

The Convergence Computing method for calculating the roots of number reduces the process down into two key instructions

- Bit shifts
- Additions

**HOWEVER!** Not only is floating point arithmetic computationally taxing, it is also not possible to shift floating point values.

Pseudo Code:

Assuming:

$M$  = input

$K$  = bits of precision

```
f = 1.0
f_sqrt = 1.0
for i = 0 to K - 1 do
     $\mu = f \cdot (1 + 2^{-i}) \cdot (1 + 2^{(-i)})$ 
     $\mu\_sqrt = f\_sqrt \cdot (1 + 2^{(-i)})$ 
    if  $\mu \leq M$  then
        f =  $\mu$ 
        f_sqrt =  $\mu\_sqrt$ 
return f_sqrt
```

# Software solution: Fixed Point Arithmetic

1. Introduce a scale factor to the floating point input and algorithm
2. Convert input into 32 bit integer
3. Perform calculation
4. Convert integer output back into type float
5. Remove the scale factor from the output

37.75 -> 4.22s !

```
[jonathonsquire@seng440 src]$ ./qRoot.exe
The Square Root of 1.000 is 1.000000
The Square Root of 2.750 is 1.400940
The Square Root of 3.000 is 1.442200
The Square Root of 4.250 is 1.619751
The Square Root of 5.000 is 1.709900
The Square Root of 6.500 is 1.866089
The Square Root of 7.999 is 1.999512
```

Increased efficiency  
by 90%!

48.05 -> 4.86s !

```
[jonathonsquire@seng440 src]$ ./qRoot.exe
The Cube Root of 1.000 is 1.000000
The Cube Root of 2.750 is 1.400940
The Cube Root of 3.000 is 1.442200
The Cube Root of 4.250 is 1.619751
The Cube Root of 5.000 is 1.709900
The Cube Root of 6.500 is 1.866089
The Cube Root of 7.999 is 1.999512
```

# Software solution: C Optimizations

## Pipelining

```
u = f+(f>>0);
u_sqrt = f_sqrt + (f_sqrt>>0);
u = u+(u>>0);

for (i=1; i<16; i++){
    if (u <= local_M){
        f = u;
        f_sqrt = u_sqrt;
    }
    u = f + (f >> i);
    u_sqrt = f_sqrt + (f_sqrt >> i);
    u = u + (u >> i);
}
if (u <= local_M){
    f = u;
    f_sqrt = u_sqrt;
}
```

## Operator Reduction of Loops

```
for(i=1;!(i&16);i++){
    ...
    if((u-local_M-1)&2147483648){
        ...
    }
}
```

## Operator Reduction of First Iteration

```
register int32_t u = f<<2;
register int32_t u_sqrt = f_sqrt<<1;
```

## Loop Unrolling

```
...
if (u <= local_M){
    f = u;
    f_sqrt = u_sqrt;
}
u = f + (f >> i);
u_sqrt = f_sqrt + (f_sqrt >> i);
u = u + (u >> i);

if (u <= local_M){
    f = u;
    f_sqrt = u_sqrt;
}
i++;
u = f + (f >> i);
u_sqrt = f_sqrt + (f_sqrt >> i);
u = u + (u >> i);
...
```

# Software solution: Assembly optimizations

arm-linux-gcc -S ...

```
str    fp, [sp, #-4]!  
add    fp, sp, #0  
sub    sp, sp, #36  
str    r0, [fp, #-8]  
mov    r1, #16384  
str    r1, [fp, #-32]  
mov    r2, #16384  
str    r2, [fp, #-28]  
ldr    r3, [fp, #-8]  
str    r3, [fp, #-24]  
ldr    r1, [fp, #-32]  
mov    r1, r1, asl #2  
str    r1, [fp, #-16]  
...
```



arm-linux-gcc -S -O3 ...

```
movle   r0, #16384  
movgt   r0, #65536  
add     r2, r0, r0, lsr #1  
add     r2, r2, r2, lsr #1  
movle   r3, #16384  
movgt   r3, #32768  
cmp     ip, r2  
movge   r0, r2  
add     r1, r0, r0, lsr #2  
add     r1, r1, r1, lsr #2  
add     r2, r3, r3, lsr #1  
movge   r3, r2  
...
```

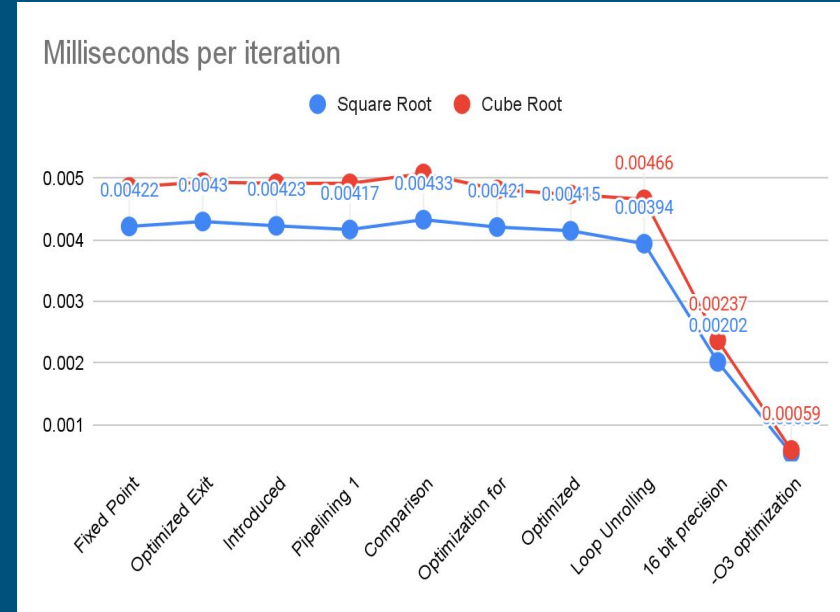


Pigeonhole optimization

```
...  
add     r3, r3, r3, lsr #1  
...
```

# Discussion

- Introducing predicates did not always result in faster runtimes
- Each optimizations impacted the different algorithms differently
- All of the load/str were replaced by add and move instructions
- Assembly file had gotten larger with optimization but was faster as the individual instructions were much less taxing





# Vertical or Horizontal Machine?

---

**Vertical Machine:** only one pipeline stage is executed at any given time

**Horizontal Machine:** set number of pipeline stages can be executing simultaneously

Considering each iteration is dependant on the previous iterations computation, horizontal machines introduce unnecessary complexity, resulting in excessive NOP instructions.

Therefore: a vertical machine is sufficient for this project

# Conclusion

---

- Took a lot of time to get fixed point working, but once working we optimized fast
- Had an increase of 8411% in performance in the cube root time

# Questions?

---